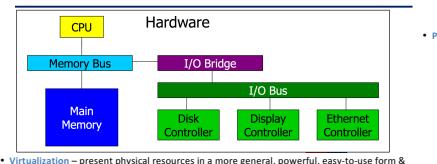
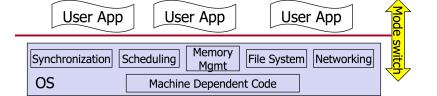


INTRODUCTION

- OS goals = 1. Convenience for user; 2. Efficient operation of computer system
- May contradict, e.g. fast windows = convenience, supercomputer = efficiency
- OS Roles: virtual machine -> interface to HW; resource allocator -> shares resources; control program -> execution of program, prevent error
- Overcome of Computer System:
 - Red line separates user software and OS software (kernel)
 - Storage: processor registers (L1, L2, L3 cache), main memory (RAM), aux memory (Disk)



Virtualization - present physical resources in a more general, powerful, easy-to-use form & present illusion of multiple/unlimited resources (goal #1 + 2)

- Limited Direct Execution - set up CPU so that next instruction is fetched from code of a process to be executed (no overhead b/c CPU in fetch-decode-execute loop)
 - Direct = process given direct CPU, (some) memory access, other resource access
 - Indirect = need to restrict operations a process can perform → trap to kernel after certain time intervals, or b/c illegal operations, system calls, etc.

Hardware Support for OS

- CPU dual-mode operation: user mode & system mode (supervisor/privileged/monitor)

- Mode bit in HW designate some instructions as **privileged instructions**
- Privileged Instruction in user mode → protection fault, trap to OS (handle safely)
- Ex. Instructions: setting mode bit; disabling interrupts; halting the CPU

- Interrupts - HW signal, causes CPU to jump to pre-defined instruction (interrupt handler)

- A boot OS file in Interrupt Descriptor Table; physical addr stored in IDT Register

- 2. CPU execution (infinite) loop: fetch inst at PC → decode instr → execute instr

- checks for interrupts at the same time

- 3. interrupt, signal from the hardware

- 4. CPU switch to interrupt handlers

- 5. Interrupt PC value stored by OldPC register

- 6. IDTR: interrupt number used to set PC to start of interrupt handler in the IDT

- 7. Execution continues

- HW Interrupts → asynchronous events

- ex. Device control signal some event (I/O complete), periodic timer interrupt

- SW Interrupts → synchronous; signal errors/requests OS service from user program

- Exceptions = raised when abnormal conditions occurred during execution

- Traps = an exception by user process -ex. syscalls

- Interrupts support OS goal #2

- 1. Illegal instructions causes software generated interrupts (**exceptions**)

- 2. Periodic HW timer interrupts ensure OS gets control at regular intervals

Bootstrapping

- HW stores BIOS (basic input-output system) in non-volatile memory (NVRAM)

- On computer startup BIOS executes, loads boot sector of disk into memory

- loads boot loader into RAM

- OS startup in hardware mode

- 1. Initialized internal data structures, create first process (**init**)

- 2. Switch to user mode, run first process, wait for something to happen

System Call

- = call that invokes the OS; for resource/service managed by OS

- 1. User program calls C library function w/ args; C library pass syscall + args to OS

- 2. Interrupt (trap) signals CPU → switch to kernel mode in OS code, invoke handler

- 3. Syscall returns, switch to user mode; jumps to next instruction in the application

- Extra level of indirection → enforce separation of user-space & kernel-space

THREADS & PROCESSES

- Processes - A program in execution, active entity; OS abstraction for execution, aka job, task

- OS manages processes by keeping track of their state (**Process State**)

- Only 1 in the Running state; others in New, Ready, Blocked, or Exit state

- Timeout = timer interrupt that shifts control back to the OS

- OS data about process in **Process Control Block (PCB)**

- Program counter, add of the next instr

- CPU scheduling info, process priority

- Memory Management info page tables

- Accounting info, resource use, etc.

- I/O status info, list of open files

- Context Switch - switch CPU to another process, pure overhead; by:

- 1. Saving the state of old process;

- 2. Load saved state for the new process

Lifecycle of Process:

- Program to Process (creation)

- 1. Source file → object file;

- 2. Object file → executable;

- 3. Create new process (AS, PCB);

- 4. Dispatch process (running)

Process Destruction:

- Process voluntarily releases all resources on **exit()** → add strace, files closed, PID & exit status retained (**zombie**) till parent cleans up → context switch

- Threads - a single control flow through a program

- Control Flow = an order in which individual instructions are executed and evaluated → we'll use it's way through the program

- Each thread has its own stack, but OS doesn't protect threads stacks in same process from each other → Fix: add an incoherent guard ring (page of memory) between stacks

- Global variables & heap allocated variables of process still shared b/w threads

- Inter-process communication need extra work/b/c diff processes = isolated

Kernel-level Threads (aka. Light-weight Processes)

- Originally a process encapsulates ownership of resources + execution state

- Modern OS thread abstraction now encapsulates execution state to make concurrency cheaper → b/c less state to allocate and init

- Management of threads: schedules threads, thread operations in kernel

- Limitation: too much overhead (thread operations still req. syscalls)

User-level Threads

- User-level threads managed by run-time system (user library)

- Small & fast: represented by PC, registers, stack, small thread control block (TCB)

- Creating threads, switching b/w threads, synchronizing threads done via procedures - no kernel involvement

- Limitation: User-level threads are invisible to the OS → OS can make poor decision

- Schedule process only if ready; de-schedule, reallocation, blocking a lock

- Hybrid Kernel & User threads - use kernel b/c user-level threads

- Associate user thread w/ kernel thread; multiplex user thread on top kernel threads

CONCURRENCY

- Synchronization - mechanism that restrict the possible interleaving of executing threads/processes → manages shared resources / purpose

- 1. Enforce single use of a shared resource (**critical section problem**)

- 2. Control order of thread execution - ex. a parent can wait for child to finish

- Note: Local Variables are not shared in a thread's own stack. Global Variables are shared (stored in static data segments); Dynamic/Heaps objects shared [malloc]

- Critical Section problem: concurrent threads manipulate a shared resource

- Given: A set of threads T_1, T_2, \dots, T_n . A set of resources shared between threads

- A segment of code which accesses the shared resources → critical section (CS)

- Thread/Process Life Cycle, dependent on scheduler/physical CPU
- There are queues for each state (ready, block, etc.) b/c easier to search for a process
- Process changes states = PC uncleared from old queue, linked to the new queue
- CPU bursts - points in threads execution where it needs the CPU to execute instructions
- I/O burst - happens in threads execution when it's waiting for a lock, waiting on I/O requests to complete (inc. thread waiting on acquiring a lock)
- CPU not needed during I/O → can schedule other processes

- Ready State (part of CPU bursts) - time spent in here depends on schedule & CPU load



- Running CPU burst - time spent in here depends on schedule & CPU load
- Release
- Time-out
- Event Occurs
- Event Wait
- Note: "Input" could be signal or unlock from another thread?

- Blocked Waiting (no starvation) - if some thread t is waiting on the CS, then there is a limit # of times other threads can enter CS before this thread

- Performance - the overhead of entering and exiting the CS is small with respect to the work being done within it

- o W/o synchronization outcome depends on the order in which the accesses b/w all threads take place b/w context switches → a race condition

- o Ex. On uniprocessor thread T1 and T2 share a variable X

- T1 increments X -> X = X+1

- Design a protocol that threads can use to cooperate w/ synchronization; that:

- 1. Each thread acquires its own lock to enter its CS, in its entry section

- 2. Spin wait or follow by lock section

- 3. Remaining code is the remainder section

- Assume no special hardware instructions, no restrictions on the # of processors

- Assume that basic machine language instructions (LOAD, STORE, etc.) are atomic

- If two such instructions are executed concurrently, the result is equivalent to their sequential execution in some unknown order

Lock algorithms

- Peterson's Algorithm (can be extended to N threads) uses shared turn and flag vars

- 1. Set own flag (indicate interest) and set turn to self

- 2. Spin wait while turn == self AND other has flag set

- If one of these conditions break - the other thread has exited CS

- If both threads enter CS same time, turn will be set to one or other not both

- my_work(id, t_id) {

- 1. id = turn; t_id = 0 or 1

- 2. flag[id] = true; // remainder section

- 3. turn = id; // entry section

- 4. while (flag[id] == turn == id) {

- 5. // busy waiting

- 6. }

- 7. }

- 8. }

- 9. flag[id] = false; // exit section

- 10. turn = id; // remainder section

- 11. }

- o Lamport's Banker Algorithm - Upon entering each thread gets a lock; lower # served next

- In case of tie, thread with the lowest id (unique & totally ordered) is served first

Scheduling Algorithms

- Non-preemptive scheduling - Once the CPU been allocated to a thread, the threads keeps the CPU until it terminates or blocks

- Preemptive scheduling - has involuntary context switch

- Algorithm 1: First Come First Serve (FCFS) is Non-Preemptive

- On uniprocessor - can disable interrupts before entering CS (prevents context switches)

- Disabling interrupts is insufficient on a multiprocessor → need special atomic instructions

- Spinlock - uses Atomic instructions: test-and-set

- Record the value of variable

- Set the variable to some non-zero value

- Return old value (hardware executes this atomically)

- Read queue is circular → each thread allowed to run for quantum (time slice)

- before being preempted and put back on the ready queue

- quantum = infinity, RR = FCFS; quantum = 0, RR = processor sharing

- Want quantum w/ the shortest expected processing time

- Problem: starvation possible if short jobs keep arriving w/ a long job in queue

- Pre-emptive variant: Shortest Remaining Time scheduling

- Algorithm 3: Round Robin (RR) is Preemptive

- Read queue is circular → each thread allowed to run for quantum (time slice)

- before being preempted and put back on the ready queue

- quantum = infinity, RR = FCFS; quantum = 0, RR = processor sharing

- Want quantum w/ the shortest expected processing time

- Problem: starvation possible if short jobs keep arriving w/ a long job in queue

- Pre-emptive scheduling is better than round robin

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

- Problem: starvation still occurs if short processes keep arriving w/ a long job in queue

- Solution: RoundRobin with quantum = 1

Translation Lookaside Buffer (TLB) cache recent PPN → PTE, lookup by VPN (16 → 64 entries)

o Part of MMU → TLB; rate loop up in parallel = 1 CPU cycle

o TLB exploit → frequency for processes to access same files over time

o Managing TLBs can be hardware-loaded or software-loaded (addressing)

o OS needs to ensure pages are consistent

o Cache miss → need to validate to make copy in TLB

o When TLB misses, new PTE has to be loaded, a cached PTE must be evicted → use TLB replacement policy (implement HW, often simple)

• Address Translation Possibilities

o Common case: process is executing on the CPU, and it issues a read to a virtual address

o Each read address goes to the TLB in the MMU (all done by hardware)

o 1. TLB does lookup using the page number of the address

o 2. Common case is that the page number matches, returning a Page Table Entry (PTE)

o for the mapping for this address (Minor Faults here)

o 3. TLB validates that the PTE protection allows for reads (Major Faults here)

o PTE marked valid & read permission is on

o 4. PTE specifies which physical frame holds the page

o 5. MMU combines physical frame & offset into a physical address

o 6. MMU performs the memory access and returns results back to CPU

Page Faults → when happens address translate is incomplete

o Minor Fault – TLB does not have a PTE mapping this virtual address; 2 possible cases:

o 1. Hardware TLB → MMU loads PTE from page table in memory

o → OS has already set up the page table so that the hardware can access it directly

o 2. Software TLB → Trap to OS, look up & load PTE from page table into TLB

o Returns from exception, retry memory access

o Now there's a PTE for the requested address; minor b/c NO U/O needed

o Major Fault – PTE exists, but memory access violates PTE protection bits; trap to OS:

o 1. R/W/X = op not permitted on page, may send fault back up to process

o 2. Invalid – virtual page not allowed/not in physical memory

o → Virtual page not allowed in address space → OS sends segmentation fault back

o Page not in memory (Major Fault here) OS allocates frame, swap it in from disk

o Minor Fault must proceed a major fault

Temporal Locality – local ref recently likely be ref again → ex. i var in loop structure;

• Spatial Locality – recently ref local ref likely to be ref soon → ex. array being iterated

• Memory Policies

o Fetch Policy – when to fetch a page – somewhat important w.r.t saving time

o Replacement Policy – how to put the page in memory – least important

o Fetch Policy: Demand Paging – pages loaded from disk when referenced; evicted to disk when memory is full and there's more demand → cause Major Page Fault miss

o Allocates a page frame, reads page from disk

o Dirty pages need to be written to disk b/c page in memory is diff from the page in disk

o → Write Eviction → indicated by Modify/Dirty bit in PTE

o Clean pages don't but, still need to know where on disk to read them from again

o Timing disk read is initiated when process needs page; but only fault on 1 page at a time

Fetch Policy: Pre-paging (Prefetching) – predict future page use at time of current fault

o Based on: Spatial & Temporal locality or history

o Good b/c disk good at large transfers; bad at individual reads

o If guess wrong – loaded unnecessary things, evicting something important (tradeoff)

Placement Policy – in paging systems, memory management hardware can translate any virtual-to-physical mapping equally well; why? pre-mapping for what's needed?

o Non-Local Policy – always map to same physical page → any processor can access entire memory, but local memory is faster for some CPUs dependent on location

Replacement Policy – how to pick a victim page to evict if memory is full

o Goal of replacement algorithm reduce the fault rate by selecting the best victim

o Best page to evict = one never used again; next best = one won't be used for longest time

o Belady's Algorithm (OPT/MIN algorithm) – lowest fault rate or any ref stream

o IDEA: replace the page that will not be used for the longest period of time

o Problem: have to know the future perfectly

o Used as a yardstick to compare w/ other page replacement algorithms

o Random replacement = lower bound for replacement algorithms

Replacement: First-In-First-Out (FIFO)

o IDEA: Maintain a list of pages in which order they were paged in; evict the oldest page

o Problem: uncertain whether 1st page might not be used again b/c no additional info.

o Belady's Anomaly – fault rate might increase when algorithm is given more memory

o Condition: 2 memory size large & small; set S: 1. At least one page in L contained within S; 2. S contains at least one page not in L

o Used as a yardstick to compare w/ other page replacement algorithms

Replacement: Least-Recently-Used (LRU)

o IDEA: evict page that hasn't been used for longest time in past

o Problem: looping pattern that's just larger than the # frames

o Exact implementation is costly, need to approximate

o Option 1: timestamp every reference, evict oldest page

o Problem: need to examine every page on eviction; find one w/ oldest time stamp

o Option 2: keep pages in a stack;

o On reference move the page to the top of the stack

o On eviction, replace page at bottom

o Problem: need costly to manipulate stack on EVERY memory reference

Replacement: Second Chance Algorithm

o IDEA: FIFO, but inspect ref bit when page is selected

o If ref bit = 0, replace the page; if ref bit = 1, clear ref bit;

o Combine w/ Modily bit to create 4 classes of pages:

o clean + never recently used; dirty + not recently; clean + recently used; dirty + recently used

o Implementation: all physical page frames in big circle (clock)

o Linux uses second-chance clock algorithm

o Replacement: Count-based replacement

o IDEA: count number of uses of a page

o Least-Frequently-Used (LFU) – replace the page used least often

o Problem: pages that are heavily used at one time sticks around; even if not needed

o Most-Frequently-Used (MFU) – favors new pages

Page Replacement – Replacement Policies assumed replacement algorithm is ran and a victim page selected when a new page needs to be brought in → too costly on every page fault

o Maintain pool of free pages, Run algorithm when pool too low → low "water mark"

o Free enough pages to at once replenish pool "high water mark"; on fault grab from list

OS161 Implementations

o Coremap – used to track the use of each physical page frame of memory

o An array of 16MB pages, entry structure, 1 page per frame

o Software-managed TLB → software & reader/writer page table

o Each region has a logical page → page table entries

Fixed vs. Variable Space Allocation – determine how much memory to allocate to processes

o Fixed space algorithms – each process is given a limit of pages it can use

o Local Replacement – When reaches the limit, it replaces from its own pages

o Variable space algorithms – process' set of pages grows/shrinks dynamically

o Global replacement – one process can run in for the rest, take up all of memory

o Local Replacement – replacement and set size are separate

Working Set Model – used to model dynamic locality of a process' memory usage

o IDEA: get a sense of how many pages this process needs

o Working Set WC(L, t) = {pages | P was ref'd in the time interval (t, t+delta)}

o t = time; delta = working set window (measured in page refs)

o Problem: how to determine delta, how do we know when working set changes → hard

o Working Set as an abstraction, utilization still needed

Page Fault Frequency (PFF) – varies space available for allocation w/ ad-hoc approach, monitor fault rate

o Threshing: OS times data I/O → time executing user programs; no useful work done

o When threshing, QoS guaranteed; Page replacement algorithms should avoid this

o Can't detect which page should be in memory due to read faults

o Solution: by more RAM

Sharing – shared memory allows processes to share data using direct memory references

o Have PTEs in both L1 and L2 cache in the same physical frame can have diff protection values

- Must update both PTEs when page becomes evicted/invalid
- Can map share memory at same/different virtual addresses in each address space
- Problem: shared pointers inside shared memory segments are invalid
- Copy-On-Write

FILE SYSTEMS

- File (management) systems – long-term (persistent) information storage, requires
- o 1. Store very large amounts of information
- o 2. Information must be terminated the process using it
- o 3. Multiple processes must be able to access info concurrently
- o Need: implement an abstraction for secondary storage (files); Organize files logically (directories); Permit sharing of data between processes, people, and machines; Protect data from unwanted access (security)

- File – a named collection of data w/ some properties (attributes)

- File Access Methods – general-purpose FS = simple methods; DB more sophisticated

- o Sequential access – read bytes one at a time, in order

- o Direct access – random access; get block by byte/number

- o Record access (DB) – fixed or variable length; Indexed access (DB)

- o File Attributes – part of FS metadata (that "other" info besides filename, file data)

- o File Operations – creating (find space in FS, add entry in dir mapping filename & attributes to location)

- o Deleting a file – Truncating (keep attributes, erase (part) of file contents)

- o Handling file operations – search for entry assoc w/ the given filename

- o Index into this table used on subsequent operations → no need to search twice

- o When files first used actively, information stored in 2 levels

- o 1. A System-Wide Open-File Table stores file's attributes info.

- o 2. A Per-Process Table records all files that each process has opened → holds current position for each process

- o Entries in the Per-Process Table points to the entry in the system-wide Open File Table b/c process independent info

- o Open File Table → System-Wide Open File Table

- o Open File Table → console/device/somefile.txt → anotherfile.txt → sample.txt

- o Directories – list of entries – i.e., file names w/ associated metadata

- o Multi-level directories (/, /usr/, /usr/local/, /home/)

- o Current Working Directory – the directory which filenames or paths are specified w.r.t.

- o Absolute path – starts from root of the directory tree

- o Directory's entries list ordered/random – sorted by algorithm that reads the directory

- o Directory Implementation – usually a tree-structure

- o Links – sharing implemented w/ dir entry called a link (pointer to another file/dir)

- o Hard Links – secondary dir entry is identical to the first

- o Soft Links – symbolic links; if it moves or renames, Hard Link follows

- o Directory entry refers to file that holds the true path to the linked file – level of indirection

- o A path is resolved upon access → if file moves/filename changes, Soft Link doesn't follow

- o Issues w/ Acyclic Graphs – a file may have multiple absolute names w/ Hard Links

- o Deletion: when can the space allocated to a shared file be deallocated, reused

- o Symbolic links: deletion of a link = OK; deletion of the file entry itself de-allocates space and leaves the link pointers dangling

- o Hard links: keep a reference count

- o Sharing: permissions of directory w/ hard link might update

- o Path Name Translation; ex. open /one/two/three in FS

- o 1. Open directory / (the root directory, well known, can always find)

- o 2. Search for one → get location of one in the directory entry

- o 3. Open directory one

- o 4. Search for two, get location of two

- o 5. Open directory two

- o 6. Search for three, get location of three

- o 7. Open the three

- o System spends a lot of time walking dir paths – this is the reason open call is separate from read/write call; OS will cache for performance

- o File System Management: FSs are designed to be fast → built by gluing together sub-systems from multiple physical partitions w/ each device (partition) stores a single file system

- o Mount point – empty directory in partition

- o File Protection/Permission – protection system dictates whether a given action ("how") performed by a given subject ("who") is allowed

- o IDEAS: who is doing what to whom

- o Types of Access

- o None – Reading, Changing Protection

- o Knowledge – Appending, Deletion

- o Execution – Updating, UNIX: R/W, R/X only

- o 2 Approaches to protection

- o Access Control List (ACL) – for each object maintain a list of subjects and their permitted actions → object-centric, easy to grant/revoke/manage

- o Problem: when objects are heavily shared, ACLs become large (UNIX user groups)

- o Capabilities – for each subject, maintain a list of objects and their permitted actions

- o Easier to transfer objects → easier to move objects

- o Problem: to revoke capabilities → keep track of all subjects that have capabilities

Objects

- o One block → location of root directory, location on disk always known

- o Free Space Bitmap – determines which blocks are free/allocated; on disk/cached in RAM

- o Remaining space = disk blocks used to store files/directories

- o Block size (usually 4KB, defined by a FS); disk space allocated w.r.t num of blocks

File System Implementation Overview

- o Master Block – determines location of root directory, location on disk always known

- o Free Space Bitmap – determines which blocks are free/allocated; on disk/cached in RAM

- o Remaining space = disk blocks used to store files/directories

- o Block size (usually 4KB, defined by a FS); disk space allocated w.r.t num of blocks

Directory Implementation

- o Option 1: Linear list of filenames + pointers to their data blocks

- o Requires linear search to find entries – easy to implement, slow to execute

- o Option 2: Hash Table; hash a filename to get pointer to entry in the linear list

Disk Layout Strategies

- o How to find all the blocks of a file? 3 options:

- o Contiguous allocation – like memory, fast, simplifies direct access

- o Problem: inflexible, cause fragmentation, need compaction → slow

- o Linked/Chained Structure – each block points to the next, the next's pointers to 1st

- o Problem: need to read sequentially, hierarchy, bad for all others

- o Indexed Allocation – good for sequential, linear, data

- o Need index Block to point to data blocks; may need multiple index blocks

- o Handles random access better, still good for sequential access

- o UNIX INODE – indexed structure for files

- o Many file metadata is stored in an INODE

- o INODE smaller than a block → need to read multiple INODEs

- o 12 and 16 byte INODEs → address of 1st 128 Blocks in the file

- o Suppose Disk Block # = address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

- o Suppose Disk Block Address = Address of Disk Block w/ address of Disk Blocks

<li